

Supplementary material for Photographic Text-to-Image Synthesis with a Hierarchically-nested Adversarial Network

Anonymous CVPR submission

Paper ID 2823

1. Training and Architecture Details

The training is similar to the standard alternative training strategy in GANs, which alternatively updates the generator and discriminators. The Adam optimizer [3] is used. The initial learning rate is set as 0.0002 and decreased by half for every 100 epochs (50 for COCO). The model is trained for 500 epochs in total (200 epochs for COCO). We configure the side outputs at 4 different scales where the feature map resolution is equal to 64^2 , 128^2 , 256^2 , and 512^2 , respectively. For the local image loss of these 4 side outputs, we set $R_1 = 1$, $R_2 = 1$, $R_3 = 5$, $R_4 = 5$. For example, R_1 refers to 64^2 . These numbers are not fine-tuned but are set empirically. We believe there exists better configurations.

All intermediate conv layers except from the specified ones in Section 3.4 use 3×3 kernels (with reflection padding). Some other normalization layers are experimented (i.e. instance normalization [5] and layer normalization [1]) since they are used by recent advances [7, 2]. But they are not satisfactory.

With respect to the generator, we use 1-repeat residual blocks for the generator till the 256^2 resolution. The input of the generator is a $1024 \times 4 \times 4$ tensor. As the feature map resolution increases by 2, the number of feature maps is downsampled also by 2, at 8, 32, 128, 256 sizes. To generate 512^2 images, we pre-train the generator to 256^2 due to the limitation of the GPU memory. We use a 3-repeat res-block followed by a stretching layer to upscale the feature map size to $32 \times 512 \times 512$. and a linear compression layer to generate images. Since the 256^2 image already captures the overall semantics and details, to encourage the 512^2 maintain these information, we use a l1 reconstruction loss to ‘self-regularize’ the generator in case the discriminator confuses the expected output.

All source code will be released for better understanding.

2. More Qualitative Results and Analysis

In this section, we demonstrate more sample results for the three datasets.

Figure 1 compares our results with StackGAN. For each input, 6 images are randomly sampled. The results demonstrate better visual quality than StackGAN. Figure 2 shows the results on the CUB bird dataset. All the outputs of a model with different resolutions are also shown. As can be observed in the two figures, our method can generate fairly vivid images with different poses, shape, background, etc. Moreover, the images with different resolutions, which are side outputs of a single model, have very consistent information. More and more image details can be observed as the resolution increases. Figure 3 shows the results on the Oxford-102 flower dataset. Very detailed petals can be generated with photographic colors and saturability.

Figure 4 shows the examples on the COCO dataset. COCO is much more challenging than the other two datasets since it contains natural images from very different scenes containing hundreds of different objects. As can be observed in the shown samples, our method generate semantically consistent and meaningful images, which correctly interpret the input text.

However, it is worth to notice that, although our method significantly improves existing methods [6, 4] on COCO, we realize that generating fine-grained details of natural scenes is still challenging. The generation of the real natural image domain remains unsolved. We expect to address this problem as the future study.

This is a tiny bird
with a large
protruding tan chest
and a short black beak
that has grey wings

HDGAN



StackGAN



Medium sized bird,
red crown to orange
fades to his tail,
orange abdomen and
belly, wing edges are
gray

HDGAN



StackGAN



This small bird has
black wings and a
yellow and black
spotted belly along
with a small, pointy
beak

HDGAN

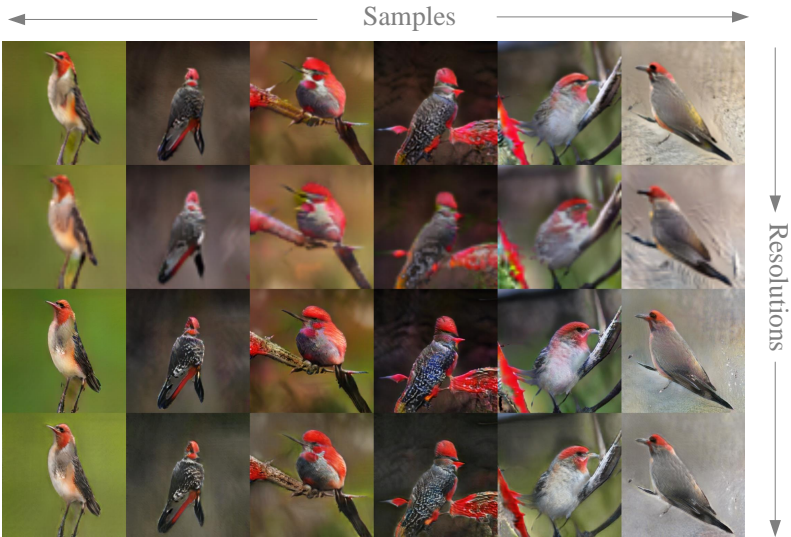


StackGAN



Figure 1. Sample results on CUB compared with StackGAN. For each input, 6 samples are shown with resolutions of 64^2 and 256^2 , respectively. As can be obviously seen, our HDGAN can generate very consistent content in images of different resolutions. Moreover, our generated images show more photographic color and contrast.

A small bird with a red crown and straight bill sits perched atop a branch



A small yellow/green bird with black wings and white wingbars




The bird has sharp pointed beak with grayish yellow throat, with white eyering



Figure 2. Sample results on CUB. For each input, 6 samples with resolutions of 64^2 , 128^2 , 256^2 , and 512^2 are shown in 4 rows, respectively.

This flower is purple and red in color, and has petals that are striped



This flower is pink and yellow in color, with petals that are skinny and oval.



Flower has petals that are pale pink with yellow stamen




Figure 3. Sample results on Oxford-102. For each input, 6 samples with resolutions of 64^2 , 128^2 , and 256^2 are shown in 3 rows, respectively.

References

- [1] J. L. Ba, J. R. Kiros, and G. E. Hinton. Layer normalization. *arXiv preprint arXiv:1607.06450*, 2016.
- [2] Q. Chen and V. Koltun. Photographic image synthesis with cascaded refinement networks. *ICCV*, 2017.

A close up of a plate of food containing broccoli



A photo of a train on a bridge above a river



A beach on a sunny day with a bunch of people on it.



A living room filled with lots of furniture



A man riding a kiteboard on top of the ocean



A very tall cathedral towering over a city



A group of people carrying ski equipment while walking on snow covered ground.



People playing with kites outside in the desert



Figure 4. Sample results on COCO. We show 8 256^2 samples in very different scenes.

[3] D. Kingma and J. Ba. Adam: A method for stochastic optimization. *arXiv preprint arXiv:1412.6980*, 2014.

[4] S. Reed, Z. Akata, X. Yan, L. Logeswaran, B. Schiele, and H. Lee. Generative adversarial text to image synthesis. *ICML*, 2016.

[5] D. Ulyanov, A. Vedaldi, and V. Lempitsky. Instance normalization: The missing ingredient for fast stylization. *arXiv preprint arXiv:1607.08022*, 2016.

[6] H. Zhang, T. Xu, H. Li, S. Zhang, X. Wang, X. Huang, and D. Metaxas. Stackgan: Text to photo-realistic image synthesis with stacked generative adversarial networks. In *ICCV*, 2017.

[7] J.-Y. Zhu, T. Park, P. Isola, and A. A. Efros. Unpaired image-to-image translation using cycle-consistent adversarial networks. *ICCV*, 2017.